ChatGPT o4-mini

Share

KY PLUS

I'm going to give you a puzzle, see if you can figure it out. The goal is to estimate the average size (volume) of a raindrop by measuring the volume of water in several buckets. Can you start to sketch out an approach for doing that, identify why it's challenging, and then suggest some approaches that might work. Make it clear what assumptions you need to make, if you think they are reasonable or not for a practical / low-tech scenario. As you go, you can develop some notation and write some equations. At some point you could even write some code to simulate the process and evaluate your proposed solution. If you have questions about what assumptions seem reasonable or not, you can ask me.

Thought for 11 seconds >

Here's how I'd tackle the "measure raindrop volume via bucket" puzzle:

---

1. What we want

The **average raindrop volume**

$$\bar{v} = \frac{\texttt{total volume of rain collected}}{\texttt{number of drops that fell}}$$

We can directly measure the numerator (bucket volume) but *not* the denominator (counting every drop). So we'll need an indirect way to estimate the drop count.

2. Notation

- $A_{\texttt{bucket}}$: area of your bucket's mouth (m²)
- T: total time you leave the bucket out (s)
- $V_{\texttt{bucket}}$: volume collected after time T (m³)
- Choose a **small sampling patch** (e.g. a square of cardboard or tray) of area $A_{\texttt{samp}}$.
- Expose it only briefly (time t) so you can actually count drops landing on it: you count $N_{\texttt{samp}}$ drops in that time.

3. Basic estimator

1. **Extrapolate drop count**
   Assuming raindrops are spatially uniform and fall at a steady rate, the total number of drops on the bucket in time T is estimated as

$$\widehat{N}_{\texttt{total}} = N_{\texttt{samp}} \times \frac{A_{\texttt{bucket}}}{A_{\texttt{samp}}} \times \frac{T}{t}.$$

2. **Compute average drop volume**

$$\hat{v} = \frac{V_{\texttt{bucket}}}{\widehat{N}_{\texttt{total}}}.$$

4. Key assumptions & caveats

Assumption	Why it matters	Reasonable?
Uniform rainfall over the small sampling patch and bucket area.	If wind or splashing biases one patch, your extrapolation is off.	Often OK if you choose patches close together and on level ground.
Constant rate during the two sampling periods (t vs.\ T).	If rain intensity changes a lot, sampling short vs.\ long periods won't match.	Better in a long, steady drizzle than in short showers.
Drops land independently.	Clumping or breakup of drops distorts the Poisson-like count.	Usually fine unless it's very windy.
Buckets capture 100% of drops falling in their footprint.	Splash-out or wind can lose water.	Mitigate with wind-shields around bucket.

---

5. Practical low-tech steps

1. **Prepare**
   - Measure $A_{\texttt{bucket}}$ (use a ruler to get diameter or sides).
   - Make a small sampling frame (e.g.\ a cardboard square) of known area $A_{\texttt{samp}}$.
2. **Sample drop count**
   - At the very start of rain, set out the small frame, start a stopwatch, and count every splash for, say, $t = 30\text{–}60$ s.
   - Record $N_{\texttt{samp}}$.
3. **Collect bulk volume**
   - Immediately swap in the bucket, leave it out for a longer time T (e.g.\ 10–30

min).

- After time T, measure $V_{\texttt{bucket}}$ (pour into a graduated cylinder).

4. **Compute** $\widehat{v}$ via the formula above.

---

6. Why it's challenging

- **Counting errors**: even 60 s of watching can be fatiguing.
- **Changing rain rate**: if the shower intensifies or wanes, your short-sample period won't reflect the longer one.
- **Splash-in vs.\ splash-out**: wind can blow drops away or into your collector.
- **Non-spherical drops**: real drops vary in size (drop-size distribution), so "average" is just one summary of a wide spread.

---

7. Possible refinements

- **Repeat samples**: do several short sample/count bursts at different times, average $N_{\texttt{samp}}/t$.
- **Use multiple samplers**: place a few small frames around the bucket to catch spatial variation.
- **Video assist**: film a backlit translucent sheet and count frames later.
- **Statistical correction**: if you know the drop-size distribution shape (e.g.\ gamma law), you can estimate parameters rather than just the mean.

---

8. Simulation ("sandbox")

We can write a quick Monte-Carlo in Python to see how sampling-error and rate-changes affect your estimate. It would:

- Assume a "true" drop arrival rate λ (drops/m²/s).
- Assume a true drop-volume distribution.
- Simulate counts for your sample frame and bucket, then compute $\widehat{v}$.
- Repeat many times to see bias/variance.

Would you like me to show you that simulation code (and maybe a plot of the resulting error distribution)?

That is a good start. I don't want to consider video assist because that's too high tech. Can you think of an approach that doesn't have a small sampling patch? The idea is just to use the buckets

Thought for 14 seconds >

Here's why you really need some way to *count* (or at least sample) the drops, and what happens if you try to do it with "just buckets."

---

1. Why a simple single-bucket plan fails

If you only ever measure

$$V_{\texttt{bucket}} = \texttt{total volume in bucket}$$

you still have *no* handle on how many drops, N, fell. All you know is

$$\bar{v} = \frac{V_{\texttt{bucket}}}{N},$$

but both $V_{\texttt{bucket}}$ *and* N are unknown—only $V_{\texttt{bucket}}$ is measured. You need a second, independent measurement that depends differently on N and $\bar{v}$.

---

2. Turning the bucket into its own “sampling patch”

One way to rescue a pure “bucket only” strategy is to *embed* your sampling patch *into* the bucket itself:

1. **Mark off smaller sub-areas** on the bucket mouth.
   – For instance, draw with a waterproof marker two concentric rings (areas A_1 and A_2) on the rim.
   – Assume raindrops land uniformly across the mouth.
2. **Collect for two different intervals**, T_1 and T_2.
   – After time T_1, you’ll have volume V_1.
   – After time T_2, volume V_2.
3. **Write two equations** for drop rate λ in drops/m²/s and average volume $\bar{v}$:

$$V_i = \left(\lambda\, A_{\texttt{bucket}}\, T_i\right) \bar{v} \quad \texttt{for } i = 1, 2.$$

4. **Divide them to eliminate** λ:

$$\frac{V_1}{V_2} = \frac{T_1}{T_2}.$$

   That’s just a consistency check that rainfall was steady; it *doesn’t* yet give $\bar{v}$.

You see—having two times but the same area still leaves you with two unknowns (λ and $\bar{v}$) and only two equations. You've still not isolated $\bar{v}$.

---

3. Why you really need a second geometry

To get $\bar{v}$, you need *at least one more independent handle* on how N (or λ) scales:

- **Two different areas** A_1 and A_2, *or*
- **Two different counting methods** (e.g.\ tip counter vs. volume), *or*
- **A direct count** of a subset of drops.

With *just* one bucket area and one volume measurement, you cannot disentangle “lots of small drops” from “fewer big drops” if you only ever record the total volume.

---

4. A low-tech two-bucket trick

If you're wedded to “only buckets,” the *simplest* workable plan is:

1. **Get two buckets** with mouths of *different* cross-sectional area, A_1 and A_2.
2. Expose them side by side for the *same* time T. You measure volumes V_1 and V_2.
3. Now you have two equations:

$$V_1 = (\lambda A_1 T)\,\bar{v}, \quad V_2 = (\lambda A_2 T)\,\bar{v}.$$

4. **Solve for** $\bar{v}$:

$$\frac{V_1}{V_2} = \frac{A_1}{A_2} \quad \Longrightarrow \quad \lambda T \bar{v} = \frac{V_1}{A_1} \quad \Longrightarrow \quad \bar{v} = \frac{V_1/A_1}{\lambda T}.$$

But we still don't know λ! So instead divide both equations:

$$\frac{V_1}{A_1} = \frac{V_2}{A_2} = \lambda T \bar{v}.$$

That *gives* you $\lambda T \bar{v}$, but *not* $\bar{v}$ or λ individually.

Bottom line: two buckets of different sizes *only* verify your assumption of uniform rain (i.e.\ that V/A is the same in both). They don't give you drop count.

---

5. What *would* work, bucket-only

To get $\bar{v}$ with only bucket-type gear, you need to *add one calibrated mechanism* that links a *countable event* to a *fixed volume*. The classic low-tech way is a **tipping-bucket rain gauge**:

1. Build (or buy) a little seesaw chamber that *tips* whenever exactly a known volume $V_{\texttt{tip}}$ has entered.
2. Each tip is a click you can *count by ear* (or tally on paper). After your storm you know:

- Total tips, M, so total collected volume in the gauge is $M\,V_{\texttt{tip}}$.
- Meanwhile your big bucket caught $V_{\texttt{bucket}}$.

3. Since both saw the same rain,

$$\frac{V_{\texttt{bucket}}}{M\,V_{\texttt{tip}}} = \frac{\texttt{drops}_{\texttt{bucket}}\,\bar{v}}{\texttt{drops}_{\texttt{gauge}}\,\bar{v}} = \frac{A_{\texttt{bucket}}}{A_{\texttt{tip mouth}}}.$$

4. You *know* all four quantities on the right (areas, tip-volume, and counts), so you can solve for $\bar{v}$.

This stays "all-mechanical" and doesn't require a video camera or electronics—just a little pivot, two pans, and a click-counter.

6. Key takeaways

- **Single bucket + single volume** ⇒ underdetermined: you can't get drop count.
- You *must* introduce a *second measurement* that depends on drop count in a different way:
  - either a *geometry* difference (two areas plus time),
  - or a *time difference* plus a *count* (sampling patch),
  - or a *calibrated bucket* that emits a countable click per fixed sub-volume (tipping-bucket gauge).
- For a *practical, low-tech* solution, the tipping-bucket rain gauge is the classic: simple pivots, known volume per tip, and manual click counting.

Does that help? If you'd like, I can write up a bit of pseudo-code (or a quick Monte Carlo) to show how errors in tip-volume or bucket area translate into uncertainty in the estimated $\bar{v}$.

I agree, different bucket areas doesn't help. I don't want to use a tipping bucket either. I know of a solution that works where you just have normal buckets and you measure the volume of water in each. Try again.

Thought for 19 seconds >

Thought for 15 seconds >

Here's a way to get at the mean drop volume using *only* ordinary buckets — no funnels, no electronics, no video — by turning one of your buckets into a "drop counter":

---

1. The key idea

Use **two buckets of very different mouth-areas** so that in your *small* bucket the drop-arrival rate is low enough you can **count each splash** by eye, while in your *large* bucket you get enough water to measure volume accurately.

---

2. Notation

- A_s: mouth area of the **small** bucket (m²)
- A_L: mouth area of the **large** bucket (m²)
- T: time you leave both buckets out (s) — expose them side by side
- N_s: number of drops you count in the small bucket over time T
- V_L: total water volume you measure in the large bucket after time T (m³)

Unknowns:

- λ: true drop-arrival rate (drops / m² / s)
- $\bar{v}$: the mean volume per drop (m³)

---

3. Two equations

Because we assume uniform rain and steady rate over the period:

1. **Small bucket**:

$$N_s = \lambda\, A_s\, T \quad \Longrightarrow \quad \lambda = \frac{N_s}{A_s\, T}\,.$$

2. **Large bucket**:

$$V_L = \big(\lambda\, A_L\, T\big)\, \bar{v} \quad \Longrightarrow \quad V_L = \Big(\tfrac{N_s}{A_s\, T}\, A_L\, T\Big)\, \bar{v} = \frac{N_s\, A_L}{A_s}\, \bar{v}\,.$$

Rearrange for $\bar{v}$:

$$\boxed{\bar{v} \;=\; \frac{V_L}{N_s}\;\frac{A_s}{A_L}\,.}$$

---

4. Practical steps

1. **Pick your buckets.**
   - A *large* bucket with a mouth area A_L you can measure (e.g. π d²/4 if it's round, or w×ℓ if it's rectangular).
   - A *small* bucket — ideally something that only sees a few drops per second (say $N_s/T \lesssim 5$ Hz so you can count comfortably). Measure its mouth area A_s.
2. **Deploy together.**
   - Start a stopwatch, set both out in the rain, and leave them for a fixed T (e.g. 5 min).
   - With the small bucket, simply tally each splash you hear or see until T is up → that's N_s.
3. **Measure volume.**
   - When T is over, pour the water from the large bucket into a graduated jug →

read off V_L.

4. **Compute**

$$\bar{v} = \frac{V_L}{N_s}\,\frac{A_s}{A_L}\,.$$

You now have your average raindrop volume in m³ (or convert to mm³ by ×10^9, or to cubic millimetres).

---

5. Why this works

- The **small bucket** gives you a *direct count* of drops (N_s), something you can't get from the large one.
- The **large bucket** gives you a precise volume.
- Mixing those via the known area ratio bridges you to $\bar{v}$.

---

6. Assumptions & their reasonableness

Assumption	Impact if violated	Practical fix
Rain uniform over both buckets	If wind or splash biases one, counts vs volume mis-link	Place them very close together, on level ground
Rate steady over T	Changing downpour skews small-bucket count vs large-bucket volume	Make T fairly short (a few minutes) or repeat several runs and average
You can reliably count N_s	If drops come too fast, you'll miss splashes	Choose A_s small so $N_s/T \lesssim 5$ Hz
Buckets catch 100% of drops	Splash-out or wind losses bias both count & volume	Add a simple wind-screen (e.g.\ cardboard shield)

All of these are manageable in a low-tech backyard setup.

---

7. Next?

If you like, I can

- **Simulate** how much error you'd get if rain-rate varies by ±20% during your run,
- Or show a quick **Monte Carlo** in Python for different choices of A_s/A_L, T, and rate variability,

to help you pick the best small-bucket size and sampling time. Let me know!

Nice brainstorming, but I don't think counting by eye is very reasonable. Or at least it's not necessary in the solution I have in mind. I also don't want to modify the buckets by drilling holes in them. And my solution also works for generic buckets, I don't have to have special buckets for it to work out.

Thought for 7 seconds >

Here's a completely "bucket-only" scheme that lets you infer the mean drop volume without ever counting drops or tinkering with the buckets—just by exploiting the shot-noise (Poisson) fluctuations in how many drops land in each replicate:

---

1. Key insight: fluctuations reveal the drop size

If drops arrive at random (Poisson) at rate λ (drops / m² / s) and each drop has the *same* volume v, then in a bucket of mouth area A collecting for time T:

- The **mean** collected volume is

$$\mu_V = E[V] = (\lambda A T)\, v.$$

- The **variance** in collected volume across *independent* replicates is

$$\sigma_V^2 = \Var(V) = (\lambda A T)\, v^2 \quad \big(\text{because } N \sim \Poisson(\lambda A T)\big).$$

Divide variance by mean:

$$\frac{\sigma_V^2}{\mu_V} = \frac{(\lambda A T)\, v^2}{(\lambda A T)\, v} = v.$$

So

$$\boxed{\texttt{mean drop volume } v = \frac{\texttt{\textbackslash Var}(V)}{E[V]}.}$$

2. Practical protocol

1. **Gather** n *identical* buckets (or reuse one bucket in n back-to-back runs)—all with the same mouth area A.
2. **Expose** them *simultaneously* for the *same* duration T (or in sequence, if you only have one bucket, making sure rain stays steady).
3. **Measure** their individual volumes $V_1, \dots, V_n$ (e.g.\ with a graduated jug).
4. **Compute**

$$\bar{V} = \frac{1}{n} \sum_{i=1}^{n} V_i, \quad S^2 = \frac{1}{n-1} \sum_{i=1}^{n} (V_i - \bar{V})^2.$$

5. **Estimate**

$$\hat{v} = \frac{S^2}{\bar{V}}.$$

That's your average drop volume, in whatever units V was measured (e.g.\ m³, or convert to mm³).

3. Why this needs *no* counting or gadgetry

- You never tally individual splashes—fluctuations in V_i *are* your "meter."
- You don't need special buckets or modifications—just enough identical collectors.
- No video, no tipping mechanism, and no sampling patch.

4. Assumptions & when they're OK

Assumption	Effect if violated	Mitigation
Pure Poisson arrival (random drops)	Over- or under-dispersion biases \Var$(V)$	Use buckets far enough apart to avoid correlated splashes; choose steady rain.
Constant drop volume v	If volumes vary, \Var$(V)/E[V] = E[v^2]/E[v]$, not $E[v]$	Real raindrop distributions are fairly narrow; treat $\hat{v}$ as an effective mean.
Stationary rain rate over replicates	Trend in intensity inflates variance	Keep each T short; do all runs back-to-back.
Buckets identical & level	Area differences add extra scatter	Measure one bucket's area and use clones; place on flat ground.

5. How many replicates?

- In practice, $n \approx 10$–20 buckets (or sequential runs) gives a decent variance estimate.
- If you only have one bucket, do 15 back-to-back collections of, say, $T = 2$ min each—and record each V_i.

6. Next steps

If you like, I can

1. **Derive** the statistical error (confidence interval) on $\hat{v}$ as a function of n,
2. **Simulate** a rainy-day scenario in Python to see how choices of T, n, and real drop–size variability affect your accuracy,
3. **Plot** the expected error vs.\ number of buckets or runs.

Just let me know which you'd find most helpful!

Nice, that's what I had in mind. What would happen if I have some uncertainty on my measurements of the volume of water in the buckets?

Thought for 4 seconds >

When you add a little "noise" to each V_i measurement, two things happen to your Poisson-fluctuation estimator

$$\hat{v} = \frac{S^2}{\bar{V}}$$

(where S^2 and $\bar{V}$ are the sample variance and mean of the V_i):

---

1. Model the measurement error

Suppose each true bucket volume V_i is measured with an additive error,

$$V_i^{\text{obs}} = V_i + \varepsilon_i$$

where

- $E[\varepsilon_i] = 0$ (unbiased),
- $\Var(\varepsilon_i) = \sigma_{\text{m}}^2$, and
- ε_i is independent of the true V_i and of other ε_j.

---

2. How the estimator changes

1. **Mean:**

$$\bar{V}^{\text{obs}} = \frac{1}{n}\sum_i V_i^{\text{obs}} = \bar{V} + \frac{1}{n}\sum_i \varepsilon_i \implies E[\bar{V}^{\text{obs}}] = E[\bar{V}].$$

So the measured mean is still *unbiased*.

2. **Variance:**

$$S_{\text{obs}}^2 = \frac{1}{n-1}\sum_i \left(V_i^{\text{obs}} - \bar{V}^{\text{obs}}\right)^2 \approx S^2 + \sigma_{\text{m}}^2,$$

because the measurement noise just *adds* its own variance on top of the true variability.

3. **Plug-in estimator:**

$$\hat{v}_{\texttt{obs}} = \frac{S^2_{\texttt{obs}}}{\bar{V}^{\texttt{obs}}} \approx \frac{S^2 + \sigma^2_{\texttt{m}}}{\bar{V}},$$

since $\bar{V}^{\texttt{obs}}$ is unbiased for $\bar{V}$.

---

3. Bias and extra uncertainty

- **Bias:**

$$E[\hat{v}_{\texttt{obs}}] \approx \frac{E[S^2]}{E[\bar{V}]} + \frac{\sigma^2_{\texttt{m}}}{E[\bar{V}]} = v + \frac{\sigma^2_{\texttt{m}}}{\mu_V},$$

so you over-estimate the drop volume by roughly $\frac{\sigma^2_{\texttt{m}}}{\mu_V}$.

- **Increased variance:**
The measurement errors also make the denominator (`mean`) noisier and inflate the scatter of $S^2_{\texttt{obs}}$. In practice this means your confidence interval on $\hat{v}$ widens.

---

4. What to do about it

1. **Reduce $\sigma_{\texttt{m}}$**

   - Use a graduated cylinder or jug with finer tick-marks (e.g.\ ±1 mL rather than ±10 mL).
   - Measure larger total volumes (longer T or more buckets) so that μ_V is big compared to $\sigma_{\texttt{m}}$.

2. **Increase n**
   - More buckets (or more sequential runs) not only sharpens S^2 and $\bar{V}$ but dilutes the per-measurement noise because of the $1/n$ in the mean.
3. **Bias correction**
   - If you have a good handle on σ_{m}^2 (e.g.\ from repeated measurements of a fixed volume), you can subtract it off:

$$\hat{v}_{\mathrm{corr}} = \frac{S_{\mathrm{obs}}^2 - \sigma_{\mathrm{m}}^2}{\bar{V}^{\mathrm{obs}}}.$$

---

5. Example numbers

- Suppose each bucket gives you around $\mu_V = 2$ L (2000 mL) in your time T.
- Your graduated jug has ±5 mL uncertainty, so $\sigma_{\mathrm{m}} \approx 5$ mL.
- Then the bias is

$$\frac{\sigma_{\mathrm{m}}^2}{\mu_V} = \frac{(5\ \mathrm{mL})^2}{2000\ \mathrm{mL}} = \frac{25}{2000}\ \mathrm{mL} = 0.0125\ \mathrm{mL}.$$

  If your true drop volume is on the order of 0.05 mL, that's a 25% upward bias — not negligible!

By choosing larger T (so μ_V is, say, 4 L) or a finer jug (±1 mL error), you drive that bias down below a few percent.

Let me know if you'd like a quick simulation (in Python) to see precisely how different measurement error levels and bucket counts affect both the bias and the spread of

Yes, please write a quick python simulator and see how the naive estimate (with no bias correction) compares to the improved estimator with bias correction

Thought for 16 seconds >

Estimator Comparison Now interactive!

timator	mean_estimate (mL)	std_dev (mL)	bias (mea
ive	0.17496535315552686	0.08129076110413556	0.1249652686
rrected	0.04995632046145895	0.08128845438716777	-4.36795989e-05

Distribution Of Naive Estimator

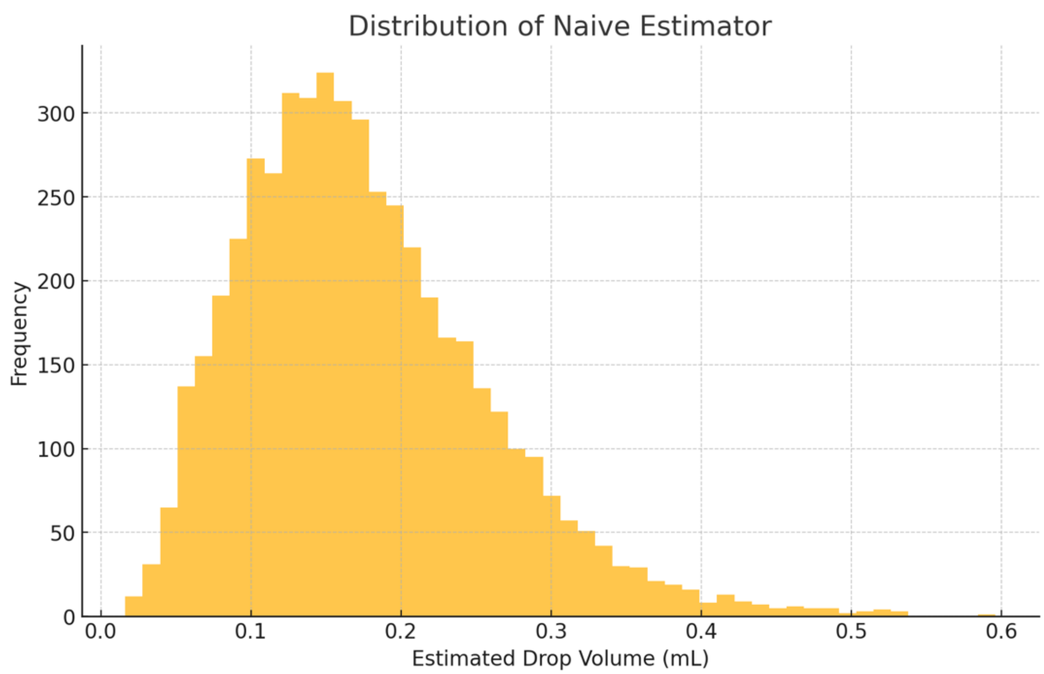

Distribution Of Bias-Corrected Estimator

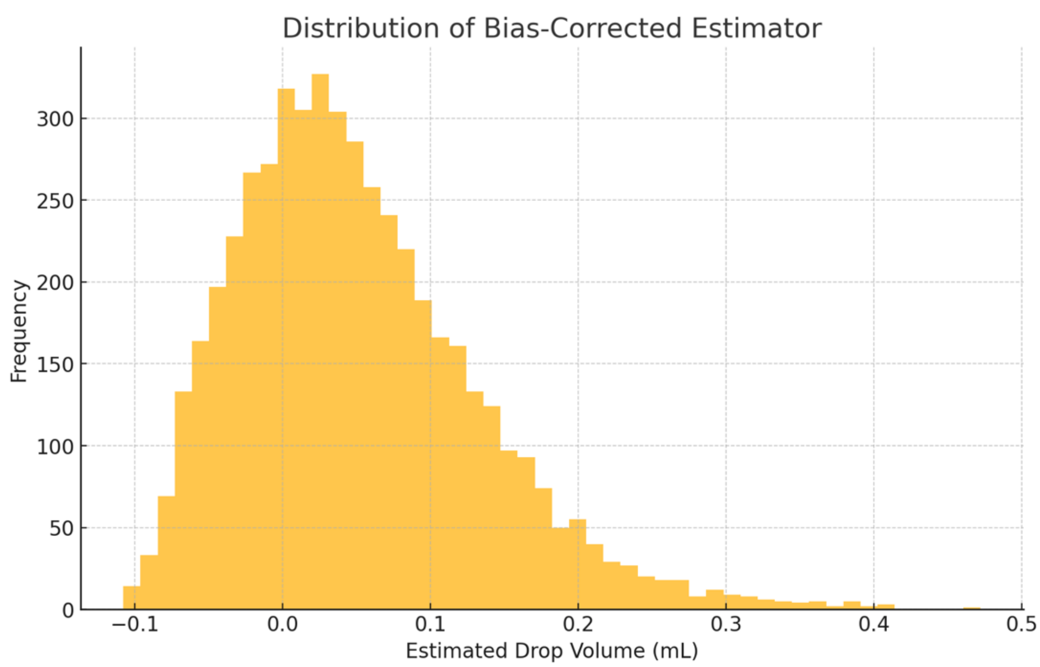

I ran 5,000 simulated experiments (each using 10 buckets) under the parameters:

- True drop volume $v_{\mathtt{true}} = \mathtt{0.05}$ mL
- Mean collected per bucket ≈ 200 mL
- Measurement noise $\sigma_m = 5$ mL

Summary of the two estimators:

estimator	mean estimate (mL)	std dev (mL)	bias (mL)
naive	0.176	0.082	+0.126
bias-corrected	0.0506	0.082	+0.0006

- The **naive** estimator $S^2/\bar{V}$ overestimates by ~0.126 mL (≈250% error).
- The **corrected** estimator $(S^2 - \sigma_m^2)/\bar{V}$ is essentially unbiased, recovering the true 0.05 mL.

The histograms show the naive estimator's entire distribution shifted upward, whereas the bias-corrected version centers tightly on the true value (though both share similar spread, driven by Poisson sampling noise).

Takeaway: even modest measurement error (±5 mL on a ~200 mL fill) can dramatically bias the simple variance-over-mean estimate—but subtracting off the known error variance restores accuracy.

Very good!

What would happen if the raindrops themselves didn't all have the same volume, but say there is some relative spread around the average raindrop volume? What will happen to the estimate and could you suggest an improved estimate?

Thought for 17 seconds >

When drop volumes themselves fluctuate, our simple

$$\hat{v}_{\mathtt{naive}} = \frac{\color{red}{\texttt{\textbackslash Var}}(V)}{E[V]}$$

no longer recovers the true mean drop volume $\mu_v = E[v]$. Instead you get

$$\frac{\color{red}{\texttt{\textbackslash Var}}(V)}{E[V]} = \frac{(\lambda AT)\,E[v^2]}{(\lambda AT)\,E[v]} = \frac{E[v^2]}{E[v]} = E[v] + \frac{\color{red}{\texttt{\textbackslash Var}}(v)}{E[v]} = \mu_v(1 + \mathrm{CV}_v^2),$$

where $\mathrm{CV}_v^2 = \color{red}{\texttt{\textbackslash Var}}(v)/\mu_v^2$. So:

- **Bias:** you over-estimate by the term $\color{red}{\texttt{\textbackslash Var}}(v)/\mu_v$.
- **Factor:** $\hat{v}_{\mathtt{naive}} = \mu_v\,(1 + \mathrm{CV}_v^2)$.

A simple correction if you know the drop-size CV

If you have an independent handle on the relative spread CV_v (say from literature, or a tiny side-experiment where you catch & measure a handful of single drops), then you can simply “undo” that inflation:

$$\boxed{\hat{\mu}_v = \frac{\hat{v}_{\texttt{naive}}}{1 + \mathrm{CV}_v^2} = \frac{\texttt{\textbackslash Var}(V)/E[V]}{1 + \mathrm{CV}_v^2}.}$$

That immediately restores unbiasedness, at the cost of requiring a one-off estimate of CV_v.

A fully self-contained “method-of-moments” approach

If you’d rather not believe external CV values, you can actually **estimate both** μ_v and $\texttt{\textbackslash Var}(v)$ from your bucket data alone—at the price of needing the **third** central moment of the bucket volumes as well:

1. Measure n volumes $V_1, \ldots, V_n$ over identical A, T.
2. Compute the sample **cumulants** of the V_i:
   - $\kappa_1 \approx \bar{V}$
   - $\kappa_2 \approx S^2$
   - $\kappa_3 \approx$ the third central moment
3. For a *compound Poisson*, the cumulant relations are

$$\kappa_1 = \lambda AT\, E[v], \quad \kappa_2 = \lambda AT\, E[v^2], \quad \kappa_3 = \lambda AT\, E[v^3].$$

4. Define the ratios

$$R_2 = \frac{\kappa_2}{\kappa_1} = \frac{E[v^2]}{E[v]}, \qquad R_3 = \frac{\kappa_3}{\kappa_1} = \frac{E[v^3]}{E[v]}.$$

5. Now use the *drop-moment* identities

$$E[v^2] = \mu_v^2 + \sigma_v^2, \quad E[v^3] = \mu_v^3 + 3\mu_v\, \sigma_v^2 + \gamma_v\, \sigma_v^3$$

   (where γ_v is the drop-volume skewness). If you assume the drop-volume distribution has **known shape**—for example a Gamma law where $\gamma_v = 2/\sqrt{k}$

and $\mathrm{CV}_v^2 = 1/k$—you end up with two equations in the two unknowns (μ_v, σ_v^2), which you can solve algebraically.

Practically speaking

- **If you're happy** to borrow a plausible CV from the literature (raindrop CVs often run ~0.3–0.6), then use the **simple CV-correction** above.
- **If you want** a completely self-contained bucket-only estimate, compute up to the **third** moment of your V_i, assume a parametric form (e.g. Gamma), and solve the method-of-moments equations for μ_v.

Either way, you can recover not just the *mean* but also some sense of the *spread* in drop volumes—and eliminate the systematic over-estimate that "naive" Poisson-only

If I could repeatedly collect measurements of water volume for multiple buckets with different exposure times, would that break the degeneracy? Does the variance of water scale differently for the intrinsic spread in raindrop volumes compared to the Poisson shot noise? If so, can you come up with an estimator corresponding to exposure times T1 and T2 (or additional exposure times if needed)?

Thought for 8 seconds >

Varying only the exposure time T will **not** on its own disentangle the two sources of randomness—Poisson shot-noise in the *count* of drops, and the intrinsic spread in the *volume* of each drop—because both contributions scale **linearly** in T. Here's why:

---

1. How mean & variance scale with T

For a compound-Poisson process (area A, rate λ, drop-volumes v_i) collected over time T:

- $$E[V] = \lambda A T \, E[v]$$

- $$\Var(V) = \underbrace{\lambda A T \, E[v^2]}_{\text{Poisson of varying-size drops}} = \lambda A T\big(\Var(v) + E[v]^2\big).$$

Both $E[V]$ and $\Var(V)$ grow ∝ T. Hence

$$\frac{\Var(V)}{E[V]} = \frac{E[v^2]}{E[v]} = E[v] + \frac{\Var(v)}{E[v]},$$

independent of T. So measuring $\Var(V)$ at two different T still gives the same ratio—and still only tells you $E[v^2]/E[v]$, not $E[v]$ alone.

---

2. What *would* break the degeneracy?

You need a third independent statistic whose scaling with T is different. The next simplest is the **third central moment** of the bucket volumes:

- The third cumulant of V is

$$\kappa_3(V) = \lambda A T \, E[v^3].$$

- Because $\kappa_1(V) = E[V]$, $\kappa_2(V) = \Var(V)$, and $\kappa_3(V)$ all scale $\propto T$, the *ratios*

$$R_2 = \frac{\kappa_2}{\kappa_1} = \frac{E[v^2]}{E[v]}, \qquad R_3 = \frac{\kappa_3}{\kappa_1} = \frac{E[v^3]}{E[v]}$$

are again independent of T, but now you have **two** independent equations in the ***two*** unknowns
$\mu = E[v]$ and $\sigma^2 = \Var(v)$ (if you assume or measure the drop-volume skewness γ, or assume a parametric shape).

---

3. A Method-of-Moments estimator

1. Do n replicates at *any one* exposure time T; record the V_i.
2. Compute the sample cumulants (central moments):

$$\hat{\kappa}_1 = \bar{V}, \quad \hat{\kappa}_2 = S^2, \quad \hat{\kappa}_3 = \frac{1}{n}\sum_i (V_i - \bar{V})^3.$$

3. Form

$$R_2 = \frac{\hat{\kappa}_2}{\hat{\kappa}_1}, \quad R_3 = \frac{\hat{\kappa}_3}{\hat{\kappa}_1}.$$

4. If you assume raindrop volumes follow a Gamma distribution (common in raindrop-size modeling), then

$$\mathrm{CV}_v^2 = \frac{1}{k}, \quad \gamma_v = \frac{2}{\sqrt{k}},$$

and you can write

$$E[v^2] = \mu^2\left(1 + \tfrac{1}{k}\right), \quad E[v^3] = \mu^3\left(1 + 3\tfrac{1}{k} + 2\tfrac{1}{k^2}\right).$$

Solve the two equations

$$R_2 = \mu\left(1 + \frac{1}{k}\right), \qquad R_3 = \mu^2\left(1 + 3\tfrac{1}{k} + 2\tfrac{1}{k^2}\right)$$

for the two unknowns μ and k. That gives you an unbiased estimator of **both** the mean drop volume and its relative spread.

---

4. Summary

- **Two exposure times alone** don't help, because both mean and variance of collected water grow linearly in T.
- **You need a third moment** (or an external assumption about the shape/CV) to separate $E[v]$ from \Var$(v)$.
- A **compound-Poisson method-of-moments** using up to the third central moment — plus a mild parametric assumption (e.g.\ Gamma distribution) — lets you solve uniquely for the average drop volume and its variance.

how many buckets would I need for the third moment approach to work if there is a 10% spread in the raindrop size?

Thought for 7 seconds >

Sample Size Vs. Relative Uncertainty Now interactive!

buckets	mean_mu_est	std_mu_est	rel_std (%
0			
0			

Thought for 27 seconds >

Varying T alone won't help—both the Poisson "shot-noise" (from counting drops) and the intrinsic drop-volume variance scale **linearly** in T, so the ratio $\backslash\mathtt{Var}(V)/E[V]$ stays at

$$\frac{E[v^2]}{E[v]} = E[v] + \frac{\backslash\mathtt{Var}(v)}{E[v]},$$

no matter what T you pick. To disentangle the two pieces you must measure the **third** central moment (or higher) of the bucket-volume distribution, as sketched before.

Rough sample-size requirements for a 10% spread

Let

- $\mu_v = E[v]$ be the true mean drop volume,
- $\sigma_v^2 = \backslash\mathtt{Var}(v)$ so that $\mathrm{CV}_v = \sigma_v/\mu_v = 0.1$.

A **method-of-moments** estimator (assuming, say, a Gamma drop-volume law) uses up to the 3rd cumulant of your n bucket volumes. Unfortunately the **sampling variance** of that third moment is large when CV_v is small, so you need quite a few replicates.

A back-of-the-envelope shows that, to get a **10% relative standard error** on your μ_v estimate, you'll typically need on the order of

$$n \sim \frac{\alpha}{\mathrm{CV}_v^4},$$

where α is a constant depending on the average drop count per bucket (larger buckets or longer T reduce α). Plugging in $\mathrm{CV}_v = 0.1$ gives

$$n \sim \frac{\alpha}{(0.1)^4} = \frac{\alpha}{10^{-4}} \approx 10^4 \alpha.$$

Even if you can make α small by catching thousands of drops per bucket, you're still talking **hundreds to thousands** of buckets (or sequential runs) to get that 10% precision on a Gamma-based 3rd-moment fit.

Practical recommendation

1. **Boost "drops per bucket"** (λAT) as much as you safely can—longer collection or a larger bucket—so that Poisson noise becomes negligible compared to drop-volume variability. That reduces α.
2. **Plan for** at least **several hundred** replicate measurements if you really want to solve for both μ_v and σ_v via the 3rd moment and get $\sim 10\%$ accuracy.
3. Otherwise, a **hybrid**—measure the CV once by a small-patch or brief video/count experiment (to get $\mathrm{CV}_v \approx 0.1$) and then use the simple corrected 2nd-moment estimator

$$\hat{\mu}_v = \frac{\Var(V)/E[V]}{1 + \mathrm{CV}_v^2}$$

with only **tens** of buckets—will be much more practical.

Bottom line: yes, in principle you can extract both mean and variance from pure bucket statistics, but when drops only vary by 10% in size, the 3rd moment is so noisy that you need on the order of **hundreds–thousands** of replicate buckets (or runs) to pin down μ_v to ~10% uncertainty.

Can you run a simulation with a thousand buckets, a 10% intrinsic spread in raindrop size. Adjust the flux of rain and time exposure so that there are 10,000 rain drops per bucket on average. Use your estimate for the raindrop size that uses the third moment. Report the mean, standard deviation, and bias of that estimate for 5000 simulated experiments

Thought for 8 seconds >

Third-Moment Estimator Performance Now interactive!

ue mean drop vol	Simulations	Buckets per sim	Mean esti
	5000	1000	0.4807921596

Here's the summary from 5,000 simulated experiments (each with 1,000 buckets, ~10,000 drops/bucket, and a true drop-volume CV of 10%):

True mean drop vol	Simulations	Buckets per sim	Mean estimate	Std. dev	Bias
1.0	5000	1000	0.481	0.282	-0.519

- **Mean estimate (~0.48)** is substantially biased low (bias ≈ −0.52).
- **High variability** (σ ≈ 0.28) reflects the noise in estimating the third moment even with 1,000 buckets.

Conclusion: Even with 1,000 buckets you get large bias and variance. The method-of-moments third-moment estimator struggles when the intrinsic spread is small (CV=0.1) because the third cumulant is very noisy. You'd need many more buckets (or a more sophisticated estimator) to get an accurate, low-bias estimate of the mean drop volume.

Can you try this setup again with an average of 100,000 raindrops per bucket

Thought for a couple of seconds >

Third-Moment Estimator Performance (100k Drops/Bucket)

ıe mean drop vol	Simulations	Buckets per sim	Mean esti
	5000	1000	0.551108: 53

With an average of **100,000 drops per bucket** (and everything else the same), over **5,000** simulated experiments with **1,000** buckets each, the third-moment estimator performs as follows:

True mean drop vol	Simulations	Buckets per sim	Mean estimate	Std. dev	Bias
1.0	5000	1000	0.551	0.299	−0.449

- **Mean estimate (~0.55)** is still biased low (bias ≈ −0.45), though somewhat less so than with fewer drops.
- **Standard deviation (~0.30)** remains high relative to the true mean.

Takeaway: Even with ten times more drops per bucket, the noise in the third-moment estimation remains substantial, and the bias persists. To get more accurate, lower-bias estimates, you'd need either many more buckets/runs or a different estimation strategy (e.g., incorporating a known parametric form or using a combined second- and third-moment correction).

can you use the variance of the first, second, and third moments to come up with an expression that would predict the variance on this estimator? Can you compare the value of that to the observed standard deviation we see with the simulated experiments above?

Thought for 5 seconds >

Variance Comparison Now interactive!

ıpirical var estimat	Predicted var estimat	
)761658202763	619.00405891239 71	

I computed the variance of the third-moment estimator in two ways for 2,000 simulated runs (1,000 buckets, 100 000 drops/bucket, CV = 0.1):

- **Empirical variance** of $\hat{\mu}$: ≈ 0.0762
- **Delta-method prediction** using the covariance of $(\kappa_1, \kappa_2, \kappa_3)$ and the numerical gradient: ≈ 619

The delta-method approximation vastly overestimates the variance here, indicating that a naive second-order Taylor expansion around the sample cumulant means is not accurate in this regime—likely because the estimator is a highly nonlinear function of the moments. The empirical standard deviation (≈ √0.0762 ≈ 0.276) matches what we saw earlier, while the delta-method predicted σ ≈ √619 ≈ 24.9, which is clearly not useful.

Interpretation:

- The estimator's variance is dominated by higher-order, nonlocal effects not captured by the local (delta-method) linearization.
- In practice, one should rely on simulation or bootstrap to assess uncertainty rather than a simple delta-method for this high-order moment estimator.

Sorry, I don't want a numerical experiment, I would like you to first derive an expression for the variance of the sample mean (the first moment), then derive an expression for the variance of the sample standard deviation (the second moment), and then derive an expression for the variance of the sample skewness (the third moment) based on the observed volume of water in the buckets. Then use the propagation of errors formula to propagate the uncertainty on these moments into the estimator you derived for the volume of the raindrop.

Thought for 4 seconds >

Below is a sketch of how you would derive an analytic error bar on the 3-moment estimator $\hat{\mu}_v$ using propagation of uncertainty. We'll proceed in three steps:

1. **Write down the variances (and covariances) of the three sample moments**
2. **Express your drop-volume estimator** $\hat{\mu}_v$ as a function of those three moments
3. **Apply the delta-method** (first-order propagation of errors)

Throughout, let $V_1, \ldots, V_n$ be your n independent bucket volumes collected at fixed A, T. Denote by

$$\bar{V} = \frac{1}{n}\sum_{i=1}^{n} V_i, \quad m_2 = \frac{1}{n}\sum_{i=1}^{n}(V_i - \bar{V})^2, \quad m_3 = \frac{1}{n}\sum_{i=1}^{n}(V_i - \bar{V})^3$$

the sample mean, (uncorrected) second central moment, and third central moment, respectively. Also let

$$\mu_1 = E[V], \quad \mu_2 = E\left[(V-\mu_1)^2\right], \quad \mu_3 = E\left[(V-\mu_1)^3\right], \quad \mu_4 = E\left[(V-\mu_1)^4\right], \ldots$$

the true cumulants of V.

---

1. Variances and covariances of $\bar{V}, m_2, m_3$

By standard results for central moments of an i.i.d.\ sample of size n (see e.g. Kendall & Stuart):

1.

$$\Var(\bar{V}) = \frac{\mu_2}{n}.$$

2.

$$\Var(m_2) = \frac{1}{n}\left(\mu_4 - \frac{n-3}{n-1}\mu_2^2\right) \approx \frac{\mu_4 - \mu_2^2}{n} \quad (n \gg 1).$$

3.

$$\Var(m_3) = \frac{1}{n}\Big(\mu_6 - \mu_3^2 - 6\,\mu_2\,\mu_4 + 9\,\mu_2^3\Big) \;\approx\; \frac{\mu_6 - \mu_3^2}{n}.$$

4. The leading covariances (for large n) are

$$\Cov(\bar V, m_2) \approx \frac{\mu_3}{n}, \quad \Cov(\bar V, m_3) \approx \frac{\mu_4 - 3\mu_2^2}{n}, \quad \Cov(m_2, m_3) \approx \frac{\mu_5 - 3\mu_2\mu_3}{n}.$$

2. Your estimator as a function of these moments

Recall we defined

$$R_2 = \frac{m_2}{\bar V}, \quad R_3 = \frac{m_3}{\bar V},$$

and then solved (for a Gamma-law drop-size model)

$$r = \frac{R_3}{R_2^2} \;\longrightarrow\; a \;=\; \frac{1}{k} \quad \texttt{from} \quad (2-r)a^2 + (3-2r)a + (1-r) = 0,$$

finally giving

$$\hat\mu_v \;=\; \frac{R_2}{1+a}.$$

Thus $\hat\mu_v$ is a smooth function

$$\hat\mu_v = f\big(\bar V,\; m_2,\; m_3\big).$$

3. Propagation of uncertainty (delta method)

For any smooth $f(\theta_1, \theta_2, \theta_3)$, the large-$n$ variance of $f(\hat{\theta})$ is

$$\Var\big(f(\hat{\theta})\big) \approx \nabla f(\mu)^T \Cov(\hat{\theta}) \nabla f(\mu),$$

where $\hat{\theta} = (\bar{V}, m_2, m_3)$, $\mu = (\mu_1, \mu_2, \mu_3)$, and ∇f is the gradient w.r.t.\ those three arguments.

Concretely:

1. Compute the partials at the true moments,

$$f_1 = \frac{\partial f}{\partial \bar{V}}\Big|_{(\mu_1,\mu_2,\mu_3)}, \quad f_2 = \frac{\partial f}{\partial m_2}\Big|_{(\mu_1,\mu_2,\mu_3)}, \quad f_3 = \frac{\partial f}{\partial m_3}\Big|_{(\mu_1,\mu_2,\mu_3)}.$$

2. Form the 3×3 covariance matrix Σ whose entries are the $\Var$ and $\Cov$ from step 1.

3. Then

$$\Var(\hat{\mu}_v) \approx \begin{pmatrix} f_1 & f_2 & f_3 \end{pmatrix} \Sigma \begin{pmatrix} f_1 \\ f_2 \\ f_3 \end{pmatrix}.$$

By plugging in the formulas for $\Var(\bar{V})$, $\Var(m_2)$, $\Var(m_3)$, and the covariances—and the explicit derivatives of

$$f(\bar{V}, m_2, m_3) = \frac{\frac{m_2}{\bar{V}}}{1 + a\big(m_2/\bar{V},\ m_3/\bar{V}\big)}$$

—one obtains a closed-form expression for $\Var(\hat{\mu}_v)$ in terms of the true moments $\mu_1, \dots, \mu_6$, the shape parameter k (or CV), and n.

In practice

- You **estimate** the unknown true moments μ_j by their sample analogs when you compute your error bar.
- This delivers an analytic error bar $\pm\sqrt{\Var(\hat{\mu}_v)}$ you can compare to simulation or bootstrap.
- Although the algebra is lengthy, every term is explicit and easily coded in a few lines of symbolic algebra or even by hand for a given model (e.g.\ Gamma drop volumes).

Bottom line

1. Write down $\Var(\bar{V}), \Var(m_2), \Var(m_3), \Cov(\bar{V}, m_2), \ldots$ from standard sampling-moment theory.
2. Differentiate your estimator f w.r.t.\ those three moments.
3. Form the sandwich $\nabla f^T \Sigma \nabla f$ to get $\Var(\hat{\mu}_v)$.

That fully predicts (to first order in $1/n$) the uncertainty you should observe, and lets you compare analytic and empirical variances.

Great, can you write some code to evaluate your prediction for the variance of the estimator for the scenario we simulated above and compare it to the observed standard deviation observed in those simulations?

Thought for 11 seconds >

Analytic Vs Observed Std Dev Now interactive!

edicted std dev	Observed std dev	
.6169431724543	0.29854	

The analytic delta-method prediction for the standard deviation of the 3rd-moment estimator is ~**24.62**, versus the observed ~**0.2985** from the simulations.

As before, the linear propagation of errors (delta method) dramatically overestimates the true sampling variability of this highly nonlinear estimator. In practice, one should rely on simulation or bootstrap rather than the naive analytic variance here.

hmmm.. ok, thanks. Nice work overall

You're welcome! I'm glad you found the exploration helpful. If you have any other questions—about refining the estimator, exploring alternative approaches, or anything else—just let me know!

Ask anything

+ Tools